

BIO392 genomic variation file formats and tools, 2021

izaskun.mallona@mls.uzh.ch

- Learning objectives
- Schedule
- Strategy
- Exam questions

Learning objectives

- ① Understand the standard NGS -> QC -> mapping -> QC -> variant calling flow
- ② Recognize a handful of the standard file formats (FASTQ, SAM, GTF, BED, VCF)
- ③ Be aware of the caveats of such formats (e.g. coordinate systems, 0/1 counting)
- ④ Manipulate standard file formats using UNIX tools
- ⑤ Run and interpret results: genomic analysis of human genetic variation due to mobile elements insertions

① Thursday 23 Sept

- Lecture on Unix, data generation, data formats, and data handling
 - Introduction to the commandline
 - Sequencing, alignment, coordinate files, variant files
- Hands-on session on UNIX and file handling
 - Working with the terminal
 - Data streams
 - Basics of awk, sed, cut
- Hands-on session on data formats
 - FASTA, FASTQ, SAM, GTF, BED, VCF

② Wednesday 29 Sept

- Hands-on session on data formats (cont)
- Small project regarding human genetic variation and interpretation
 - Retrotransposition events in 1000 genomes data

- Included within the lectures slides

Tasks/materials at a glance

- Slides
- Exercises <https://github.com/compbiozurich/UZH-BI0392/blob/imallona/course-material/2021/imallona/exercises.md>
- Papers/resources, most of which are 'extra' reads (linked within the slides and/or exercises)
- Literature/background for running the project on mobile elements insertions
 - <https://www.nature.com/articles/nmeth.1906> chromHMM paper
 - <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC2832824/>, bedtools paper
 - <https://www.nature.com/articles/nature15393>, 1000 genomes project

The project

- Background
 - Retrotransposition events introduce genetic variation in humans
- Biological question
 - How often do retrotransposons jump inside a genome?
 - Do we detect any insertion patterns? Do transposons jump inside genes? in promoters? in heterochromatin?
 - (and any other questions you decide)
- Procedure
 - Integration of 1000 Genomes Project and functional genomics data (e.g. chromatin segmentation data)
- Running the analysis
 - Either locally (Mac), [renku](#), or in your computer (easier with Linux)

- All analysis are run on open data, with free software running on a terminal = console
- To be run on a Linux, a Mac, or (a docker image in) renku <http://renkulab.io/>

- Commands/options are in typewriter font
- URLs are highlighted in blue

Exercise: Web browsing the genome

- Launch the UCSC Genome Browser
- Specify Human Assembly hg19
- Click go
- Notice the many data layers, including variation data

By default, the Genome Browser will render a genomic window with many data layers on it. How are these data encoded?

- Click UCSC Genes from the Genes and Gene Predictions section under the main genomic window.
- Click View table schema opens [knownGene table schema](#)

Schema for UCSC Genes - UCSC Genes (RefSeq, GenBank, CCDS, Rfam, tRNAs & Comparative Genomics)

Database: hg19 **Primary Table:** knownGene **Row Count:** 82,960 **Data last updated:** 2013-06-14

Format description: Transcript from default gene set in UCSC browser

field	example	SQL type	info	description
name	uc001aaa.3	varchar(255)	values	Name of gene
chrom	chr1	varchar(255)	values	Reference sequence chromosome or scaffold
strand	+	char(1)	values	+ or - for strand
txStart	11873	int(10) unsigned	range	Transcription start position (or end position for minus strand item)
txEnd	14409	int(10) unsigned	range	Transcription end position (or start position for minus strand item)
cdsStart	11873	int(10) unsigned	range	Coding region start (or end position if for minus strand item)
cdsEnd	11873	int(10) unsigned	range	Coding region end (or start position if for minus strand item)
exonCount	3	int(10) unsigned	range	Number of exons
exonStarts	11873,12612,13220,	longblob		Exon start positions (or end positions for minus strand item)
exonEnds	12227,12721,14409,	longblob		Exon end positions (or start positions for minus strand item)
proteinID		varchar(40)	values	UniProt display ID, UniProt accession, or RefSeq protein ID
alignID	uc001aaa.3	varchar(255)	values	Unique identifier (GENCODE transcript ID for GENCODE Basic)

So they are database entries with **chrom**, **start** and **end** features. This is the most standard data representation in genomics: data referring to genomic coordinates. Why?

- How big is the human genome, in both nucleotides and file sizes?

- How big is the human genome, in both nucleotides and file sizes?
 - 3 billion nt
 - 0.8 GB (the primary assembly), e.g. [ensembl data repository](#)
- How would you store information, including sequence, related to human genomes?
 - Genes
 - Variants
 - Promotes
 - Physical properties, e.g. local G+C content
 - (etc)


- How big is the human genome, in both nucleotides and file sizes?
 - 3 billion nt
 - 0.8 GB (the primary assembly), e.g. [ensembl data repository](#)
- How would you store information, including sequence, related to human genomes?
 - Genes
 - Variants
 - Promotes
 - Physical properties, e.g. local G+C content
 - (etc)
- Using file formats based in coordinates (chromosome, start, end), and not sequences

Commonly used formats


- Reference genomes
- Fasta and FastQ (Unaligned sequences)
- SAM/BAM (Alignments)
- BED (Genomic ranges)
- GFF/GTF (Gene annotation)
- BEDgraphs (Genomic ranges)
- Wiggle files, BEDgraphs and BigWigs (Genomic scores).
- Indexed BEDgraphs/Wiggles
- VCFs (variants)

- Reference genomes describe the 'consensus' DNA sequence
- (The human genome from whom? What does consensus mean?)
- Human variation aside, multiple assemblies have been released (i.e. hg18, hg19, hg38...)

Sanger sequencing Nature 409, 863 (2001)



Hierarchical shotgun Nature 409, 863 (2001)



GRCh stands for 'Genome Reference Consortium'

- Human [GRCh37](#) (hg19)
- Human [GRCh38](#)
- Mouse [mm10](#)
- Mouse [GRCm38](#)
- Zebrafish, chicken and others: <https://www.ncbi.nlm.nih.gov/grc>The Genome Reference consortium

Activity: sequence retrieval

- Retrieve the sequence of the human mitochondrial genome
- Where (databases)? how? which format?
- Does the human mitochondrial genome exist?
- If succeeded to download one, is it mouse clicking automatable/reproducible?

- Using a Web browser to retrieve genomic sequences is not efficient nor reproducible: programmatic alternatives exist
- Need of standardizing data analysis using reproducible workflows
 - Scripts for data retrieval (in bioinformatics often R or python)
 - Data storage: standards (fasta, fastq, sam, vcf...)
 - Keeping track of data analysis steps and avoiding manual editing

- What do we mean by data science reproducibility?
- In data science: avoid manual steps of data analysis using scripts plus version control systems
- Spreadsheet editors used with sequences of mouse clicks are not reproducible

- Paper: <https://www.nature.com/news/1-500-scientists-lift-the-lid-on-reproducibility-1.19970>

How could we increase reproducibility?

- In data analysis: keeping track of all steps (using scripts)
- Using data standards
- What if we don't know how to program?
- Still, switching to command-line tools and keeping track of the commands used
- Using control version systems

The terminal

- Simple command line interface
- Present in MacOS and GNU/Linux computers
- Interprets the Unix shell language (commonly bash)

- Efficient
- Scalable
- Portable
- Open


According to Peter H. Salus in A Quarter-Century of Unix (1994):

- Write programs that do one thing and do it well
- Write programs to work together
- Write programs to handle text streams, because that is a universal interface

Why in bioinformatics?

- We interpret DNA, proteins as text; Unix is for text streams
- Data are big (millions of lines of text, easily a couple of GB); spreadsheet software (Excel) cannot handle them
- We need to keep track of our analysis for the sake of reproducibility: bash scripts

Opening a terminal in MacOS



The shell (Unix shell)

- The Unix shell allows to save the sequential commands in a reproducible manner
- Activity: run the tutorial by the Swiss Institute of Bioinformatics
https://edu.sib.swiss/pluginfile.php/2878/mod_resource/content/4/couselab-html/content.html
- (Normally requires a full day)


A quick reminder on computer files

- Files are data representations stored in computers as arrays of bytes.
- File type is defined by its bytes and not by the filename extension.
- Files contain metadata.
- Importantly, plain text files are composed by bytes mapped directly to ASCII characters.
- Text editors (notepad, gedit, vim...) allow editing plain text files.
- (text files can be read without proprietary software)


Setting up the Mac text editor to save text

- Create new file
- Go to Format and select Make Plain Text
- For saving, go to File, Save As and Plain Text Encoding setting: Unicode (UTF-8).

Avoiding RTF: plain text



Avoiding RTF: plain text



Turning off autospell check

- UNIX commands are case sensitive
- By default, TextEdit capitalizes/spell checks contents
- Exercise: disable this feature
- <http://osxdaily.com/2014/05/06/turn-off-autocorrect-pages-textedit-mac/>

Organizing a shell script

- Write on top the shebang
- Write the date and what's the script about, your name and date
- Tip: comment lines start with `#`
- Tip: multiline chunks are split by a `\`
- Introduce the commands (one line each)

- Save the file (script) somewhere with a back-up system
- You could run the script to run the commands in batch typing `bash name_of_the_script.sh`

- UNIX solves the reproducibility, scalability and openness for data (text) streams, but extra software might be needed
- The importance of software versioning for reproducibility: keeping track of the software installed
- Using open source software (no blackboxes!)
- Installs can be run command-line, so specific versions can be stored and included into the analysis script

- Use the arrow left/right keys to move back/forth within the same line
- Use the arrow up/down keys to retrieve previous commands
- Use the 'tab' key to autocomplete
- Lines starting with # are comments - not evaluated
- Text can be written after the prompt (dollar sign \$)

Basic Unix commands: **man**ual

`man man` # prints the manual for the `man` command

`man sed` # commonly used editing program

exit pressing 'q'

Basic Unix commands: **p**rint **w**orking **d**irectory

```
$ pwd
```

```
# prints my path, e.g. /home/guest
```

Basic Unix commands: **c**hange **d**irectory

```
$ cd .. # goes to parent
```

```
$ cd /tmp # goes to a folder named '/tmp'
```

Basic Unix commands: **l**ist **d**irectory

```
$ ls # list files in a given path
```

```
$ ls -l # lists files and their attributes
```

Basic Unix commands: **make** **directory**

```
$ mkdir newdir # creates the directory newdir
```

Basic Unix commands: **remove** **dir**ectory

```
# removes the directory newdir (created before)
$ rmdir newdir
```

Basic Unix commands: **copy** file

```
$ cp source-file destination-file
```

```
# let's create a file using echo
```

```
#   and redirecting the echo to a file
```

```
$ echo 'lorem ipsum' > test_file
```

```
# cp test_file test_file_another
```

```
$ ls
```

Basic Unix commands: **less** printing files one screen at a time

```
$ less test_file_another
```

```
# exit pressing q
```


Basic Unix commands: check the **head** of a file

```
$ head test_file_another
```

Other common commands

cat - printing files onto the screen

mv - moving and renaming files

rm - removing files and directories

chmod - changing access permissions

grep - searching for strings in files

wc - counting words

sed - stream editing files

awk - a full language to process texts

Standard output, input, pipes, redirects

- Commands not only send results to the screen: can be redirected
- Commands to redirect streams
 - $a > b$ sends the std output of 'a' to 'b'
 - $a|b$ pipe, sends the std output of 'a' to 'b' (with buffering)
 - $a < b$ sends the std output of 'b' to 'a'

Standard output, input, pipes, redirects

```
# editing the standard ouput using pipes
echo "hello there"
echo "hello there" | sed "s/hello/hi/"
echo "hello there" | sed "s/hello/hi/" | sed "s/there/world/"
```

```
# redirecting the standard output to a file
echo hello > newfile.txt
ls -l newfile.txt
cat newfile.txt
```

```
# piping commands and redirecting to a file
echo hello | sed 's/hello/hElLo/' > newfile2.txt
ls -l newfile2.txt
cat newfile2.txt
```

Software installs - compiling bedtools (for next week)

- Some tools we use are not standard Unix programs, but domain-specific software
- BEDtools is one of the most used tools for handling coordinate-based file formats
- If interested, read paper
<https://academic.oup.com/bioinformatics/article/26/6/841/244688>

- We will download the open source code of bedtools and then (try to) generate the executables
- That is **compiling** the source, and generates binaries fit to our hardware and OS
- Also, makes the same code runnable in many platforms
- That's transparent for us: the developer wrote a Makefile that automates this
- Extra reading:
https://www.wired.com/2010/02/Compile_Software_From_Source_Code/

Compiling bedtools

(The code is available at the exercises file)

```
cd # to your home directory or wherever you decide
```

```
cd soft # the folder was created before (for kent utils)
```

```
curl -L https://github.com/arq5x/bedtools2/releases/download/v2.25.0/bedtools-2.25.0.tar.gz  
> bedtools-2.25.0.tar.gz
```

```
tar zxvf bedtools-2.25.0.tar.gz
```

```
cd bedtools2
```

```
make
```

```
alias bedtools='./bin/bedtools'
```

- Download some data/install bedtools
(<https://github.com/compbiozurich/UZH-BI0392/blob/imallona/course-material/2021/imallona/exercises.md>, section **Set up**)
- Run exercises 1-4 (same url above)
- Run the introductory course on Unix at https://edu.sib.swiss/pluginfile.php/2878/mod_resource/content/4/couselab-html/content.html

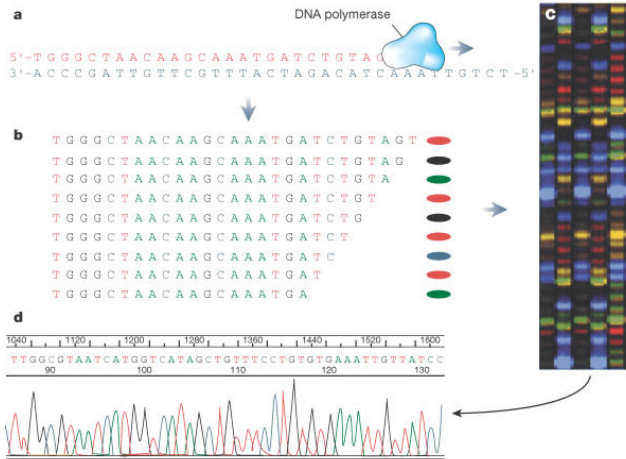
Genomic data formats

Commonly used formats


- Reference genomes
- Fasta and FastQ (Unaligned sequences)
- SAM/BAM (Alignments)
- BED (Genomic ranges)
- GFF/GTF (Gene annotation)
- Wiggle files, BEDgraphs (Genomic scores).
- VCFs (variants)

- Reference genomes describe the 'consensus' DNA sequence
- (The genome of who? who was chosen to be the 'reference?')
- Human variation aside, multiple assemblies have been released due to technical challenges/errors

Sanger sequencing Nature 409, 863 (2001)



Hierarchical shotgun Nature 409, 863 (2001)



GRCh stands for 'Genome Reference Consortium'

- Human [GRCh37](#) (hg19)
- Human [GRCh38](#) (hg38)
- Mouse [mm10](#)
- Mouse [GRCm38](#) (mm11)
- Zebrafish, chicken and others: [The Genome Reference consortium](#)

Reference genomes: FASTA

- A reference genome is a collection of contigs/scaffolds
- A contig is a stretch of DNA sequence encoded as A,G,C,T,N.
- Typically comes in FASTA format.
- ">" line contains the scaffold name
- Following lines contain the sequence (single line, 80 nt-column sized...)

Reference genomes: FASTA format

```
>NC_009902.1 Babesia bovis T2Bo mitochondrion, complete genome
TTTAAAAAAGTGTTAAAAACTTTATACATTAAAAAATTTAAACAAGTGATCATGTATAAAGTACACTTGT
TACTGTGTAAATATCAAAAACAATTTAATTTCAAAATTTTTGAAATATGTTTTTTGTGTTGTGTTATAAA
GTTTTTTTTTCAAAATTATATATGTTTGCATTTGCTGGATATAGTTCGGTCTCTGCAAACCATAAAGTCAT
CGGTATATCCTACATATGGCTTTTCATATTGGTTTGGAGTTATTGGATTTTATATGAGTATTTTGATAAGA
ACAGAATTGAGTATGAGTGGTTTAAAGATTATGACAATGGATACTCTTGAGATATACAATATGATGTTTT
(...)
```


Patches, alternate loci and primary assembly

- Primary assembly: the best known assembly of a haploid genome
 - Chromosome assembly
 - Unlocalized sequence (associated to a chromosome but whose order/orientation is unknown)
 - Unplaced sequence (not linked to any chromosome)
- Alternate loci: An alternate representation of a locus (usually highly polymorphic regions, such as the MHC region)
- Patches: A contig sequence that is released outside of the full assembly release
 - Fix: error correction
 - Novel: new sequences that will be included into the next full assembly release

Browsing genomic patches

- Activity: notice the assembly patches, i.e.
<https://www.ncbi.nlm.nih.gov/grc/human>
- Activity: browse an individual patch
http://ftp.ncbi.nlm.nih.gov/genomes/all/GCA/000/001/405/GCA_000001405.27_GRCh38.p12/README_patch_release.txt

- This we can do because the genome consortia and the science community release free data, software/toolsets
- We benefit from the Unix-like operating systems
- But the access to the same data and tools is not enough: the need for data standards

- **Fasta and FastQ**
- SAM/BAM (Alignments)
- BED (Genomic ranges)
- GFF/GTF (Gene annotation)
- BEDgraphs (Genomic ranges)
- Wiggle files, BEDgraphs and BigWigs (Genomic scores).
- Indexed BEDgraphs/Wiggles
- VCFs (variants)

Short reads sequencing

- Sequencing very short reads (50 to 150 nucleotides) is common practice
- We get hundreds of millions of short reads for each experiment
- Instead of assembling them, we map them into a reference genome
- Sequencers provide sequence and error rates assessment: fasta format is not suitable, but fastq is

FASTQ: Short read sequencing

- Next step to FASTA: including quality data
- Standard de facto for short read, high-throughput sequencing instruments such (i.e. Illumina)


```
@SRR001666.1 071112_SLXA-EAS1_s_7:5:1:817:345 length=36
GGGTGATGGCCGCTGCCGATGGCGTCAAATCCCACC
+
IIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIII9IG9IC
```

- Sequence quality is represented using Phred scores
- The sequencing quality score of a given base Q is defined by as
- $Q = -10 \log_{10} P$

Phred quality scores are logarithmically linked to error probabilities

Phred Quality Score	Probability of incorrect base call	Base call accuracy
10	1 in 10	90%
20	1 in 100	99%
30	1 in 1000	99.9%
40	1 in 10,000	99.99%
50	1 in 100,000	99.999%
60	1 in 1,000,000	99.9999%

phred scores (old school Sanger electrophoretogram)



Phred scores encoding (Wikipedia)



```
S - Sanger           Phred+33, raw reads typically (0, 40)
X - Solexa           Solexa+64, raw reads typically (-5, 40)
I - Illumina 1.3+    Phred+64, raw reads typically (0, 40)
J - Illumina 1.5+    Phred+64, raw reads typically (3, 41)
                    with 0=unused, 1=unused, 2=Read Segment Quality Control Indicator (bold)
                    (Note: See discussion above).
L - Illumina 1.8+    Phred+33, raw reads typically (0, 41)
```

Unaligned sequences (from sequencers): FASTQ

- FASTQs stands for FASTA with Qualities
- Plain text files with chunks of four lines:
 - @ identifier line
 - Sequence
 - "+" (sometimes the sequence name, again)
 - Quality scores (different encodings exist)

Example FASTQ entry

```
@SRR001666.1 071112_SLXA-EAS1_s_7:5:1:817:345 length=36
GGGTGATGGCCGCTGCCGATGGCGTCAAATCCCACC
+
IIIIIIIIIIIIIIIIIIIIIIIIIIIIIIII9IG9IC
```

Example FASTQ entry (II)

@SRR001666.1 071112_SLXA-EAS1_s_7:5:1:817:345 length=36

GGGTGATGGCCGCTGCCGATGGCGTCAAATCCCACC

+

IIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIII9IG9IC

@SRR001666.1 071112_SLXA-EAS1_s_7:5:1:817:346 length=36

GGGAGAAGGCCGCAGCCGAAGGCGACAAAACCCACC

+

IIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIII9I9999

How could you extract the sequence of each FASTQ record?

- It's always on the second line of each stanza
- How could we describe this mathematically?

How could you extract the sequence of each FASTQ record?

- It's always on the second line of each stanza
- How could we describe this mathematically?
- Use the **modulo operator** on the line number: extract lines, 2 and 6
 - $2 \bmod 4 = 2$
 - $6 \bmod 4 = 2$
- Tip: in `awk`, line numbers can be printed with the `NR` variable: `awk 'print NR,$0' filename`

- Activity: FASTQ/A exercises (exercises 5 to 14)

Working with fastq files

```
## retrieving an example fastq file
curl https://molb7621.github.io/workshop/_downloads/SP1.fq \
  > file.fastq

## transforming into fasta
awk 'NR % 4 == 1 {print ">"$1};
     NR % 4 == 2 {print}' file.fasta
```

awk: Counting the number of items in a fastq

So in fastq each data chunk is stored in four different lines. We'll need to be able to extract the first, second, third of fourth line for each block of four lines. Using awk,

```
awk 'END{print NR/4}' file.fastq
```

- NR gives the number of records (line numbers)
- FASTQ are chunks of 4 lines for each sequence
- NR/4 at the END of the file indicates the number of sequences

- These tasks can be done in many programming languages
- In bioinformatics, we normally use AWK
- <https://en.wikipedia.org/wiki/AWK>
- <http://www.grymoire.com/Unix/Awk.html>

Working with fastq files

```
## retrieving an example fasta file
curl https://molb7621.github.io/workshop/_downloads/SP1.fq \
  > file.fastq

## counting number of reads
awk 'END{print NR/4}' file.fastq

## transforming into fasta
awk 'NR%4==1{a=substr($0,2);} NR%4==2 {print ">"a"\n"$0}' \
  file.fastq
```

awk: fastq to fasta

```
awk 'NR%4==1{a=substr($0,2);} NR%4==2 {print ">"a"\n"$0}' \
file.fastq
```

- % is a modulo operator
- `NR%4==1` will retrieve the first line of a fastq chunk (header)
- `NR%4==2` will retrieve the second line (the sequence)
- the id line will be prepended with the `>` and reduced to a substring (chopped)
- This will be applied to all lines!

Still need to align the FASTQ reads to the reference genome

- Discussion: how to get rid of the sequences and to have a smaller data representation?
- Trying to transform sequence to reference genome coordinates (= aligning to the genome/mapping)
- i.e. transforming ACGCACGCACGCACGCCCC to human genome hg19
'chr10:10010-10030'

- SAM - Sequence Alignment Map.
- The standard stores where the reads (i.e. the ones we had as FASTQs) map in the reference genome

What is an alignment?

- Sequence alignment: arrange a set of sequences to identify regions of similarity/identity
- Mapping short reads against a reference genome: aligning large amounts short reads to a reference genome

Local alignments vs global alignment

- Alignment/mapping consist on finding homologies of query sequences to the reference
 - Local: aligns a substring of the query to a substring of the reference
 - Global: contains all letters from both query and reference

(reference: top, query: bottom)

1	2	3	4	5	6	7	8	9	10	11
C	G	T	C	C	G	A	A	G	T	G
			.							
★	★	T	A	C	G	A	A	★	★	★

(a) Global alignment

3	4	5	6	7	8
T	C	C	G	A	A
	.				
T	A	C	G	A	A

(b) Local alignment


1	2	3	4	5	6	7	8
C	G	T	C	C	G	A	A
			.				
★	★	T	A	C	G	A	A

(c) Semi-global alignment

Alachiotis et al, 2013

- Chromosome
- Locus (coordinate)
- CIGAR string, i.e.
 - 30M1D2M - 30 bases continuously match, 1 deletion from reference, 2 base match
- Some flags
(<https://broadinstitute.github.io/picard/explain-flags.html>)

Next generation sequencing to SAM



Pavlopoulos et al 2013

Next generation sequencing to SAM

starting
symbol

sequence
identifier

sequence

quality
score

sequence end
start QS

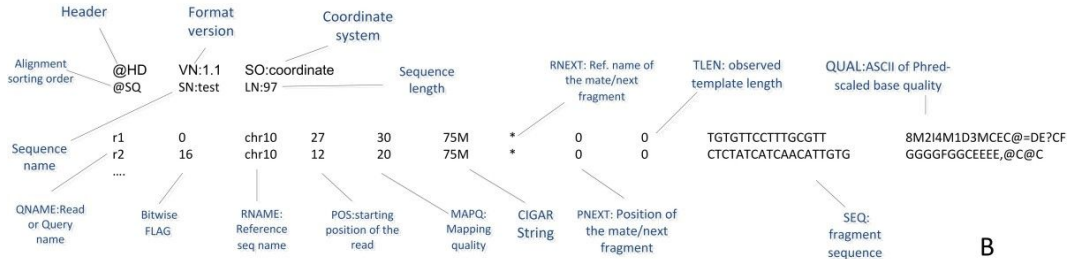
```
@HWI-EAS3X_10102_2_120_19829_1823#0/2  
TCTAACTCTTACTTAGCATAGCTGTAAAATTTTGAGTT  
+(optionally the same identifier)  
DEAEE:B:BE5EEEEED=:DEA:-AE5DDBDFFEDEEDFAE
```

Pavlopoulos et al 2013

Next generation sequencing to SAM

Coordinates 123456789...
 Reference AAATGAATAATCTCTATCATCAACATTGTGTTCCCTTTCGCTTTTAACCTTTCCT
 Reads r1 TGTGTTCCCTTTCGCTTT
 r2 CTCTATCATCAACATTGTG

A



B

Pavlopoulos et al 2013

- Extra activity: read the SAM format specification
- <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC2723002/>

- Exercise number 15


- SAM (Sequence Alignment Map) files are human-readable text files
- BAM (Binary Alignment Map) files are their binary (and compressed) equivalent

Manipulating SAM files: from SAM to BED


- Next step after getting the SAM files: extract features from reads, by coordinates and get rid of sequences
 - RNA-seq: expression levels, per gene
 - ChIP-seq: binding, per bin
 - Variant calling: variation detection
- Why? they are smaller and easier to handle
- What do we mean by feature quantification?

Manipulating SAM files: from SAM to BED

Sample A Reads



Sample B Reads



- Discussion: how to handle expression data, i.e. transcripts without introns etc?
how do we count them?
- Food for thought: <https://bedtools.readthedocs.io/en/latest/content/tools/genomecov.html>

Keep it simple: count and transform into BED files

- BED (Browser Extensible Data) files come in different flavours
- BED3: 3 tab separated columns, chromosome (scaffold), start, end
- BED6: BED3 plus name, score, strand

Just chromosome, start, end

```
chr22 1000 5000
```

```
chr22 2000 6000
```

BED3 and name, score, strand

```
chr22 1000 5000 cloneA 960 +  
chr22 2000 6000 cloneB 900 -
```

BED6 and thickStart, thickEnd, itemRgb, blockCount, blockSizes, blockStarts

Details: <https://genome.ucsc.edu/FAQ/FAQformat.html#format1>

```
chr22 1000 5000 cloneA 960 + 1000 5000 0 2 567,488, 0,3512
```


```
chr22 2000 6000 cloneB 900 - 2000 6000 0 2 433,399, 0,3601
```

How do we specify coordinates? counting from 0 or 1

- Even though BED files are standard how to count nucleotides is not
- 0-start vs. 1-start : Does counting start at 0 or 1?
- Fully-open, fully-closed, or hybrid-intervals (e.g., half-open)?
- Extra activity: browse <http://genome.ucsc.edu/blog/the-ucsc-genome-browser-coordinate-counting-systems/>

Counting: ends


Given an interval $---a-----b---$



FULLY CLOSED (HINT: THINK EN-CLOSED!)

CLOSED-START, CLOSED-END


BOTH ENDPOINTS "A" AND "B" ARE INCLUDED



FULLY OPEN

OPEN-START, OPEN-END

BOTH ENDPOINTS "A" AND "B" ARE EXCLUDED



HALF-OPEN

CLOSED-START, OPEN-END

ENDPOINT "A" IS INCLUDED, "B" IS EXCLUDED

Counting: 0s or 1s

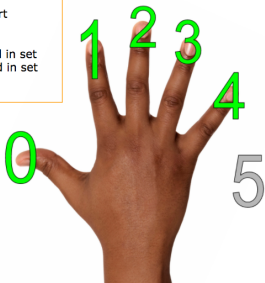
0-Start, half-open
closed-start (included)
open-end (excluded)

Range= 0-5

Size = Stop - Start
5 = 5 - 0

Start (0) **included** in set
Stop (5) **excluded** in set

COORDINATES STORED
IN THE
UCSC GENOME BROWSER TABLES



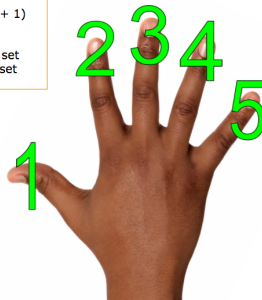
1-Start, Fully-Closed
closed-start (included)
closed-end (included)

Range= 1-5

Size = Stop - Start (+ 1)
5 = 5 - 1 (+1)



Start (1) included in set
Stop (5) included in set

COORDINATES POSITIONED
WITHIN THE UCSC GENOME
BROWSER WEB INTERFACE



Unlike the coordinate system used by other standards (GFF), BED format is zero-based half-open

Impact of different interval definitions



V Buffalo (2015), Bioinformatics data skills

How to transform 'counting systems'?

```
# generate a bed3 record
echo -e 'chr1\t0\t1000' > test.bed

# check the original file
cat test.bed

# offset all starts +1
awk '{OFS=FS="\t"; print $1,$2+1,$3}' test.bed
```

- Exercises 16 to 24

The need for further data formats

- So to sum up until now generally we have a reference genome, reads that were retrieved as FASTQ files, mapped and transformed to SAM files
- (Which step was the last one with actual sequences on them?)
- So, at last, we can answer questions
- Which fraction of the human genome is covered by exons?
- Genomic locations of SNPs associated with prostate cancer?
- Are gene bodies more variable (in terms of SNPs) than intergenic regions?
- For this we need further data analysis tools (i.e. BEDtools) on BED files and other data formats (annotations and variation)

What does annotation mean?

- Genomic annotations are layers to genomic coordinates specifying their nature



Henrik Lantz, BILS/SciLifeLab

How to store genomic annotations? GFF3

<u>Segid</u>	source	type	start	end	score	strand	phase	attributes
Chr1	Snap	gene	234	3657	.	+	.	ID=gene1; Name=Snap1;
Chr1	Snap	mRNA	234	3657	.	+	.	ID=gene1.m1; Parent=gene1;
Chr1	Snap	exon	234	1543	.	+	.	ID=gene1.m1.exon1; Parent=gene1.m1;
Chr1	Snap	CDS	577	1543	.	+	0	ID=gene1.m1.CDS; Parent=gene1.m1;
Chr1	Snap	exon	1822	2674	.	+	.	ID=gene1.m1.exon2; Parent=gene1.m1;
Chr1	Snap	CDS	1822	2674	.	+	2	ID=gene1.m1.CDS; Parent=gene1.m1;
		start_ codon						Alias, note, ontology_term ...
		stop_ codon						

Henrik Lantz, BILS/SciLifeLab

How to store genomic annotations? GTF

<u>Segid</u>	source	type	start	end	score	strand	phase	attributes
Chr1	Snap	exon	234	1543	.	+	.	gene_id "gene1"; transcript_id "transcript1";
Chr1	Snap	CDS	577	1543	.	+	0	gene_id "gene1"; transcript_id "transcript1";
Chr1	Snap	exon	1822	2674	.	+	.	gene_id "gene1"; transcript_id "transcript1";
Chr1	Snap	CDS	1822	2674	.	+	2	gene_id "gene1"; transcript_id "transcript1";
		start_ codon						
		stop_ codon						

Henrik Lantz, BILS/SciLifeLab

Why so complex? Open reading frames



Steven M. Carr

Why so complex? CDS, exons, introns, stop codons

- How many transcript does a gene have?
- Genes have transcripts, exons, CDS, start codons
- Some exons belong to several transcripts
- In GTF and GFF, parent information is needed to link different records (lines)
 - i.e. linking a transcript to its gene

- Extra documentation:

<https://www.ensembl.org/info/website/upload/gff.html> and
<https://github.com/The-Sequence-Ontology/Specifications/blob/master/gff3.md>

- Run exercises 25 to 27

Are there other relevant coordinate file formats?

- Strand-less score information
- BEDgraph (another flavour of BED)
- Wiggle (bigWig, if compressed), (even more) compact data representations

- To display continuous-valued data in track format.
- Useful for probability scores

chromA	chromStartA	chromEndA	dataValueA
chromB	chromStartB	chromEndB	dataValueB

chr19	49303800	49304100	0.50
chr19	49304100	49304400	0.75
chr19	49304400	49304700	1.00

Which are the differences between BEDgraphs and BED?

- BED, BED6, BED12?
- Advantages: the coordinates are specified, so sparsity is allowed
- Next step in file formats: trying to cover all the genome (that is, no sparsity anymore)
- Example: does it make sense to generate a BED file with GC content? ([GC content at Wikipedia](#))
- How can we store features with definite start and ends but for which the value is the primary purpose, but not their starts and ends?

- To display continuous-value data
- GC percent, probability scores, and transcriptome data.
- Data is not sparse! Wiggle data elements must be equally sized (step)

Let's construct a Wig file: a score per nucleotide

- Basic Wig file: specifies the chromosome (once), the nucleotide, and the score.
- Any improvements?

```
variableStep chrom=chr2  
300701 12.5  
300702 12.5  
300703 12.5  
300704 12.5  
300705 12.5
```

- Improved Wig: add the span if the score applies to several nucleotides

```
variableStep chrom=chr2 span=5  
300701 12.5
```

Wig with fixedStep and span

- Improved Wig: how to represent the scores of 300 nucleotides from chr3, starting from 400601, where the first 100 nt are scored 11, the next 100 nt 22, and the last 33?

```
fixedStep chrom=chr3 start=400601 step=100 span=5
```

```
11
```

```
22
```

```
33
```

- Read more on Wig files at
<https://genome.ucsc.edu/goldenpath/help/wiggle.html>

- Standard file format for storing variation data
- Unambiguous, scalable and flexible
- 8 columns:
 - 1 CHROM
 - 2 POS
 - 3 ID
 - 4 REF
 - 5 ALT
 - 6 QUAL
 - 7 FILTER
 - 8 INFO

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO	FORMAT	NA19909
11	5248232	rs334	T	A	100	PASS	AA=T ;AC=1;AF=0.0273562;AFR_AF=0.0998;AMR_AF=0.0072;AN=2;DP=22876;EAS_AF=0;EUR_AF=0;EX_TARGET;NS=2504;SAS_AF=0;VT=SNP	GT	0 1

EMBL/EBI training

Quality values: which one?

- Phred-scaled quality score for the assertion made in ALT. i.e. $Q = -10 \log_{10} P$ being $P(\text{call in ALT is wrong})$
- Read quality
- Mapping quality
- Variant calling quality

Variant calling: the VCF format

- Activity: read <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3137218/>

On coordinates, 0s vs 1s and open and closed intervals

- Remember the unusual 0 and 1 counting when manipulating GTF, GFF, BED, Wiggle and/or when using genomic positions at the UCSC genome browser
- Resource: <http://genome.ucsc.edu/blog/the-ucsc-genome-browser-coordinate-counting-systems/>

Counting in 0/1 spaces

- BED 0-based
- GTF 1-based
- GFF 1-based
- SAM 1-based, BAM 0-based
- VCF 1-based, BCF 0-based
- Wiggle 1-based
- GenomicRanges 1-based
- BLAST 1-based

V Buffalo (2015), Bioinformatics data skills

- Complete the remaining exercises
- Brainstorm; design an alternative for a common task: retrieve the human gene DEXF promoter

Retrieving fasta sequences manually (UCSC)

- Try to retrieve the DIEXF gene promoter in a non-programmatic way
- (What is a promoter? do we know where they are?)
- Choose a reference assembly, i.e. [hg38](#)
- Query gene symbol (i.e. DIEXF)
- Click into the gene (encode track)
- Click into the sequence and links item
- Specify your promoter definition

Manually downloading the DEXF promoter

← → ↻ https://genome-euro.ucsc.edu/cgi-bin/hgGene?hgg_gene=uc001hhr.3&hg ☆

Genomes Genome Browser Tools Mirrors Downloads My Data H

Human Gene DEXF (ENST00000491415.6) Description and Page Index

Description: Homo sapiens digestive organ expansion factor homolog (zebrafish) (DEXF), mRNA
Gencode Transcript: ENST00000491415.6
Gencode Gene: ENSG00000117597.17
Transcript (Including UTRs)
Position: hg38 chr1:209,828,007-209,857,565 **Size:** 29,559 **Total Exon Count:** 12 **Strand:** +
Coding Region
Position: hg38 chr1:209,828,064-209,851,447 **Size:** 23,384 **Coding Exon Count:** 12

Page Index	Sequence and Links	UniProtKB Comments	CTD	RNA-Seq Express
RNA Structure	Protein Structure	Other Species	GO Annotations	mRNA Description
Other Names	Methods			

Data last updated: 2016-03-28

[-] Sequence and Links to Tools and Databases

Genomic Sequence (chr1:209,828,007-209,857,565)			mRNA (may differ from genome)		Protein
Gene Sorter	Genome Browser	Other Species FASTA	Gene interactions	Table Schema	BioGP
CGAP	Ensembl	Entrez Gene	ExonPrimer	GeneCards	Gepis
HGNC	HPRD	Lynx	MGI	MOPED	neXtP
PubMed	Reactome	Stanford SOURCE	UniProtKB		

Manually downloading the DIEXF promoter

[Home](#) [Genomes](#) [Genome Browser](#) [Tools](#) [Mirrors](#) [Downloads](#) [My Data](#)

Genomic Sequence Near Gene

Get Genomic Sequence Near Gene

Note: if you would prefer to get DNA for more than one feature of this track at a time, try the [Table Browser](#) using the output format sequence.

Sequence Retrieval Region Options:

- ☒ Promoter/Upstream by bases
- ☐ 5' UTR Exons
- ☐ CDS Exons
- ☐ 3' UTR Exons
- ☐ Introns
- ☐ Downstream by bases
- ☒ One FASTA record per gene.
- ☐ One FASTA record per region (exon, intron, etc.) with extra bases upstream (5') and extra downstream (3')
- ☐ Split UTR and CDS parts of an exon into separate FASTA records

Note: if a feature is close to the beginning or end of a chromosome and upstream/downstream bases are added, they may be truncated in order to avoid extending past the edge of the chromosome.

Manually downloading the DIEXF promoter

```
>hg38_knownGene_uc001hhr.3 range=chr1:209827007-209827008
gtttctgctgttggttaaattggggaatgctggaacagatttgtttgctggg
actcttccaatactttcagaaaatgcgagaatagggtaggggtgggaatc
tcagacttggtgggccccatgattgatataacacacacaggcggcagaccc
taatgggtaaaagcatgtggttgcatcagttaagggttttctctcttctc
ttgctagcgtgttatcttttcttttcttttcttttcttttcttttcttgcg
agatggagtctagcttttgtcgcccaggctggagtaggctggagtgcagt
ggagtgatctcggctcattgcaacctccacctcccgggttccagcgattc
tcctgcctcacctcctgagtagctgggattacaggcgcccgcctaccacgc
ccggctgatttttgtacttttagtagagacgggggtttcaccatgtttggc
catgctgggtctcgaactcctgacctcagggtgatccgcccattctcggcctc
ccaaagtgttgagattacaggcgtagccaccgcgcccggccgctagcgt
gttatcttttctaagcatcagtttcttatctgcaacaccaggcttatta
acaagacctatctgtacactgttggtgatgaagtgagatgttcaggca
cccttaaattgttggttgatattttattgcagtatactgtaaagtcactg
cattcgactatctccgctactacacatttacgcagactgatttcataac
caaaacacaagcacaagctcatgccccgactcacgcaacccgggaagc
ccagctgcccacgttctagggctctgagaacactagtgaacgaactcccg
tgctttcaaagagctgcggtagggggcagaaccgggaaccggatgttcta
agcctgtcgtagcgagcgacgtaaaagcggtatctgctttatggcaccttg
ctttcgccgtaaagcgcagtcagcgagccacgtgcttggttgactgga
```

How to specify the DIEXF promoter in a reproducible manner?

Same concept that when we did manually:

- 1 Download the reference genome sequence (if we didn't have that already)
- 2 Download a file with all the genes (transcripts) locations (not sequences, but their coordinates)
- 3 Then select the gene we are looking for (DIEXF)
- 4 Decide what a promoter is (i.e. 2 kb upstream of the gene) and update the coordinates accordingly
- 5 Then use a specific tool to slice the full genome to only report the DIEXF promoter

How to specify the DIEXF promoter? standards

Same concept that when we did manually + some standards

- 1 Download the reference genome sequence (fasta)
- 2 Download a file with all the genes (transcripts) locations (GTF)
- 3 Then select the gene we are looking for (DIEXF) (grep/awk)
- 4 Decide what a promoter is (i.e. 2 kb upstream of the gene) and update the coordinates accordingly (BEDfile, bedtools/awk to subtract the start coordinates in a strand-aware manner)
- 5 Then use a specific tool to slice the full genome to only report the DIEXF promoter (bedtools)

- The 1000 Genomes Project profiles genetic variation in humans, including SNVs and structural variation
- Practical use case: explore the events of insertion of mobile elements using the CEU population's VCF
 - How frequent are these events?
 - Where do they happen? e.g. in gene bodies, in promoters, in heterochromatic regions?
 - Is there any biological explanation for that?

Exam questions

```
chr22 1000 5000 cloneA 960 +  
chr22 2000 6000 cloneB 900 -
```

- These records belong to a:

- ① Wiggle file
- ② VCF file
- ③ GTF file
- ④ BED6 file

```
fixedStep chrom=chr3 start=400601 step=100 span=5
```

```
11
```

```
22
```

```
33
```

- These records belong to a:

- 1 Wiggle file
- 2 VCF file
- 3 GTF file
- 4 BED6 file

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO
20	14370	rs6054257	G	A	29	PASS	NS=3;DP=14;AF=0.5;DB;H2
20	17330	.	T	A	3	q10	NS=3;DP=11;AF=0.017

- These records belong to a:

- 1 Wiggly file
- 2 VCF file
- 3 GTF file
- 4 BED6 file

```
@MN00537:51:000H2K25G:1:11101:2213:1092 1:N:0:9
CTCCAGTCCTTACTCCCATATCTAACCTCTTACCCCTACNTCATAGGTANACATTTTAATGAAT
+
FFFFFFFFFFFFFFAFFFFFFFFF=FFFFAFFFFFFFFF/AFFF#FFFFFFFFF#FFFFFFFFFFFFFFF
```

- This record belongs to a:
 - 1 Wiggle file
 - 2 VCF file
 - 3 FASTQ file
 - 4 BED6 file

Assuming a 0-based, half-open coordinate system and a chromosome named chr1 whose start is 'AACCGGTT...'. Which sequence is encoded as a BED3 record 'chr1 0 2'?

- ① AA
- ② AAC
- ③ AC
- ④ A